

Introduction to Data Science · Lab Project

Diabetes Risk Prediction

*Predicting diabetes from 21 health indicators
using supervised machine learning*

GROUP

Ali Raza · 2540010
M. Ammar · 2540004
Taha Ali · 2540008

INSTRUCTOR

Sir Zaki
Semester 2 · 2026

RESOURCES

[GitHub repository](#) · [Live Streamlit demo](#)

Abstract

This project builds a machine-learning system that predicts whether a patient is at risk of diabetes based on 21 health indicators collected through the CDC's Behavioral Risk Factor Surveillance System (BRFSS) 2015 survey. After cleaning the dataset of 70,692 records and removing 1,635 duplicates, we trained and compared four supervised classifiers — Logistic Regression, Decision Tree, K-Nearest Neighbors and Random Forest — on an 80/20 stratified train/test split. The Random Forest classifier achieved the best overall performance with 74.36% accuracy, 0.727 precision, 0.794 recall, and an F1 score of 0.759. The model identified five dominant risk factors — GenHlth (0.191), HighBP (0.166), BMI (0.149), Age (0.110), HighChol (0.079) — which align with established medical literature. The trained model is deployed as an interactive Streamlit web application that accepts user input and returns an instant diabetes-risk prediction along with the confidence score.

Keywords: Diabetes, Random Forest, Supervised Learning, Binary Classification, Feature Importance, BRFSS, Streamlit, Healthcare.

Table of Contents

1. Introduction	
2. Dataset Description	
3. Methodology & Block Diagram	
4. Data Preprocessing	
5. Exploratory Data Analysis	
6. Machine Learning Approach	
7. Model Comparison	
8. Conclusion & Future Work	

1. Introduction

Diabetes is one of the fastest-growing chronic diseases worldwide. According to the International Diabetes Federation (IDF) Atlas 2021, an estimated **537 million** adults live with diabetes, and roughly half of all cases are undiagnosed. Early screening is critical because lifestyle changes and medication can dramatically reduce complications such as heart disease, stroke, kidney failure and vision loss.

Traditional diagnostic tests (HbA1c, fasting glucose) are accurate but require a laboratory, a clinic visit and a trained phlebotomist — they are not always available in low-resource or rural settings. In contrast, the BRFSS survey collects 21 self-reported indicators (BMI, blood pressure, cholesterol, general health, age, etc.) via a phone interview. These indicators are **cheap to collect, anonymous, and available at scale**, making them ideal inputs for a machine-learning screening tool.

Project objectives. This project has two goals:

1. **Primary:** Build a supervised classifier that predicts whether a patient is diabetic given the 21 BRFSS indicators.
2. **Secondary:** Identify the most important risk factors from the data so they can be communicated back to doctors and patients.

Headline result. The chosen model — a Random Forest — achieves 74.36% accuracy and 0.759 F1 score on a held-out test set of 13,812 patients. The five dominant risk factors are: GenHlth (0.191), HighBP (0.166), BMI (0.149), Age (0.110), HighChol (0.079).

2. Dataset Description

Source. CDC Behavioral Risk Factor Surveillance System (BRFSS) 2015 — an annual telephone survey conducted across the United States. We use the cleaned, balanced version published on Kaggle by Alex Teboul.

Size & shape. Original file has **70,692 rows × 22 columns**. After dropping exact-duplicate rows we end up with **69,057 rows × 22 columns** (removed **1,635** duplicates).

Target. The column Diabetes_binary takes the value **0** (no diabetes) or **1** (diabetic / pre-diabetic). The cleaned dataset is pre-balanced: about 50 % of patients belong to each class.

Features (21). 21 health indicators organised in 3 categories — **Vital / Clinical**, **Lifestyle**, and **Self-reported / Demographics** — all listed on the next page.

First 5 rows of the dataset (df.head()):

HighBP	HighChol	BMI	GenHlth	Age	HeartDiseaseorAttack	DiffWalk	Income	Education	Diabetes_binary
1.0	0.0	26.0	3.0	4.0	0.0	0.0	8.0	6.0	0.0
1.0	1.0	26.0	3.0	12.0	0.0	0.0	8.0	6.0	0.0
0.0	0.0	26.0	1.0	13.0	0.0	0.0	8.0	6.0	0.0
1.0	1.0	28.0	3.0	11.0	0.0	0.0	8.0	6.0	0.0
0.0	0.0	29.0	2.0	8.0	0.0	0.0	8.0	5.0	0.0

Table 1: First 5 rows of the BRFSS 2015 dataset — 10 of the 21 features shown (the omitted 11 columns have identical 0/1 structure). Last column is the target. Full dataset has 22 columns and 69,057 rows.

2.1 References & Resources

2.01 Dataset

CDC Behavioral Risk Factor Surveillance System (BRFSS) 2015. Cleaned version published on Kaggle by Alex Teboul: **Diabetes Health Indicators Dataset**, 2019. Available at <https://www.kaggle.com/datasets/alexteboul/diabetes-health-indicators-dataset>.

2.02 Libraries

- **pandas** — data manipulation
- **numpy** — numerical operations
- **matplotlib**, **seaborn** — charts and visualisations
- **scikit-learn** — Random Forest, train/test split, metrics
- **joblib** — model persistence
- **streamlit** — interactive web application
- **python-pptx** — automatic PPT generation
- **reportlab** — PDF report generation

2.03 Project resources

GitHub repository: <https://github.com/muhammadammarff-design/diabetes-risk-prediction>

Live Streamlit demo:

<https://diabetesriskpredictions.streamlit.app/>

3. Methodology & Block Diagram

We follow the standard CRISP-DM data-mining workflow. The diagram below shows the three layers of the project: **Input** (the raw dataset), **Methodology** (the 4 sequential steps we perform) and **Output** (the final model results).




Figure 1: Block diagram of the project — Input (top), Methodology (middle), Output (bottom).

3.1 Top-10 features most correlated with diabetes

Rank	Feature	Correlation
1	GenHlth	0.397
2	HighBP	0.372
3	BMI	0.286
4	HighChol	0.281
5	Age	0.275
6	DiffWalk	0.267
7	Income	0.213
8	HeartDiseaseorAttack	0.207
9	PhysHlth	0.207
10	Education	0.159

Table 4: Top-10 features by absolute correlation with Diabetes_binary.

4. Data Preprocessing

Before training the model we performed the following steps. All operations are encoded in `scripts/train_model.py` (see code below).

4.1 Steps performed

- **Checked missing values:** 0 missing values found — no imputation needed.
- **Removed duplicates:** Removed 1,635 duplicate rows.
- **Converted dtypes:** All 21 feature columns cast to int (except BMI which is float).
- **Verified class balance:** Still $\approx$ 50/50 between diabetic and non-diabetic patients.
- **Final shape:** 69,057 rows $\times$ 22 columns.

4.2 Code (key lines from `train_model.py`)

Load the CSV:

```
# Path to the dataset
# DATA_PATH = "data/diabetes_binary_5050split_..."
df = pd.read_csv(DATA_PATH) # load the file into memory
df.shape                   # → (70693, 22) — rows × columns
```

Clean — drop duplicate rows:

```
# df.duplicated() returns True for every row that has appeared before
df = df.drop_duplicates().reset_index(drop=True)
# → 69 057 rows remain (1 636 removed)
```

Why drop duplicates? If we keep them, the same patient can appear in both the training and the test set. The model would 'memorise' the answer for those patients, and our reported accuracy would be artificially inflated. This is called **data leakage** — a viva-favourite concept.

5. Exploratory Data Analysis (EDA)

EDA answers two questions before we train anything: (1) Is the dataset balanced? (2) Which features visually separate diabetic from non-diabetic patients? We generated four key plots.

5.1 Class balance




Figure 2: After cleaning the dataset has 33,960 non-diabetic and 35,097 diabetic patients — almost perfectly balanced, so we don't need resampling.

5.2 BMI distribution by diabetes status




Figure 3: Diabetic patients (orange) have visibly higher BMI than non-diabetic patients (teal) — the peak shifts from 25-27 to 30-33.

6. Machine Learning Approach

6.1 Task type

- **Learning type:** *Supervised* — the target column is already labeled (0 / 1).
- **Task:** *Binary classification* — predict one of two classes.
- **Output:** Class label (0 or 1) **plus** a confidence probability.

6.2 Models compared

Logistic Regression	Fits a probability line between the two classes.	Simple and interpretable, but assumes linear decision boundary.
Decision Tree	Splits the data using yes/no questions.	Easy to interpret, but prone to over-fitting on its own.
K-Nearest Neighbors (KNN)	Predicts by majority vote among the k closest training points.	Intuitive, slow on large datasets, sensitive to feature scaling.
Random Forest (selected ✓)	Combines 200 decision trees; majority vote wins.	Robust to over-fitting, handles non-linear patterns, gives feature importance for free.

Table 5: The 4 candidate classifiers (highlighted row = our final choice).

6.3 Train / test split

We split the cleaned dataset **80 % train / 20 % test**, stratified by the target so both subsets keep the 50 / 50 class balance. Splitting is done **before** any scaling or model training so the test set is never seen during fitting — this prevents **data leakage**.

Code:

```
# from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.20, stratify=y, random_state=42)
# → train: 55 245 rows test: 13 812 rows
```

7. Model Comparison

All four models were trained on the same 80 % train split and evaluated on the same 20 % test split. The table and chart below show accuracy, precision, recall and F1 score.




Figure 7: Accuracy (teal) vs F1 score (amber) for the four classifiers.

7.1 Full metrics table

Model	Accuracy	Precision	Recall	F1
Logistic Regression	0.745	0.738	0.773	0.755
Decision Tree	0.732	0.718	0.778	0.747
K-Nearest Neighbors	0.735	0.718	0.789	0.752
Random Forest	0.744	0.727	0.794	0.759

Table 6: Side-by-side metrics for the 4 models.

Why Random Forest wins. Logistic Regression has the highest raw accuracy (74.50%) but its F1 is lower. Random Forest combines high accuracy (74.36%) with the best F1 score (0.759) and gives **feature importance for free** — a must for medical interpretation.

8. Conclusion

8.1 Summary of results

This project successfully built a Random Forest classifier that predicts diabetes risk from 21 BRFSS health indicators. Key results on the held-out test set:

Metric	Value	Meaning
Accuracy	74.36%	correct predictions
Precision	0.727	predicted diabetics correct
Recall	0.794	real diabetics caught
F1 Score	0.759	best of 4 tested models

Table 8: Headline metrics of the deployed Random Forest model.

8.2 Limitations

- **Self-reported survey data is noisy.** Patients may mis-report weight, smoking or exercise habits.
- **No lab values.** The model does not see HbA1c or fasting glucose — the gold-standard diabetes markers.
- **Balanced training data $\neq$ real-world prevalence.** The actual diabetes rate in the population is ~14 %, not 50 %; production usage would need threshold calibration.
- **Class 1 mixes diabetic and pre-diabetic patients** — they are not distinguishable in this dataset.